

Praxis: Fitting power-laws to data

In this assignment we will be combining what we learned this week to fit a power-law to some data and then visualize both the power-law and the fit.

(1) Getting the data

First we need to generate some data to fit to. We will use the degree distribution from the Barabasi-Albert graph because we know that it is power-law distributed. Create a BA graph using networkx with $n = 2500$, $m = 8$, and $\text{seed}=1$ (adjust n or m if your computer has difficulty creating the graph). The seed will prevent the graph from changing each time you run the code.

Then build a list/array of the node degrees.

In [1]:

```
# your code here
import networkx as nx
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
```

In [2]:

```
BA_graph = nx.barabasi_albert_graph(n=1200, m=8, seed=1)
```

In [3]:

```
node_degree = {node: BA_graph.degree(node) for node in BA_graph}
degree = list(node_degree.values())
```

In [4]:

```
print (degree[:10])
```

```
[83, 50, 69, 86, 111, 35, 48, 84, 133, 183]
```

(2) Find the power-law cut-off

It is normally the case in power-law distributed data that there is a value below which the power-law relation does not hold. We will express this value as x_{\min} . In order to fit the scaling exponent, we first need to estimate x_{\min} and then discard all values in the distribution below it.

You can use Aaron Clauset's [powerlaw](http://tuvalu.santafe.edu/~aaronc/powerlaws/) (<http://tuvalu.santafe.edu/~aaronc/powerlaws/>) package to determine x_{\min} . The package can be installed using pip from the commandline: `pip install powerlaw`

The package should then be available for import:

```
import powerlaw
```

And can be called using (see [documentation for additional arguments](http://arxiv.org/pdf/1305.0215v3.pdf) (<http://arxiv.org/pdf/1305.0215v3.pdf>)):

```
fit = powerlaw.Fit(some_data)
print(fit.xmin)
```

Alternatively, you can get a rough estimate of the cut-off yourself by visualizing the power-law.

In [5]:

```
# Find xmin and adjust data (your code here)
import powerlaw

fit = powerlaw.Fit(degree)
print(fit.xmin)
```

9.0

Calculating best minimal value for power law fit

C:\Users\Bao\Anaconda3\lib\site-packages\powerlaw.py:692: RuntimeWarning: in
valid value encountered in true_divide
(Theoretical_CDF * (1 - Theoretical_CDF))

(3) Find scaling exponent

With the data available and xmin estimated we can now estimate the scaling exponent for our degree distribution. Use the powerlaw package to estimate the scaling exponent.

In [6]:

```
# Your code here
fit = powerlaw.Fit(degree, xmin=9.0)
print(fit.power_law.alpha)
```

2.99321191036

C:\Users\Bao\Anaconda3\lib\site-packages\powerlaw.py:692: RuntimeWarning: in
valid value encountered in true_divide
(Theoretical_CDF * (1 - Theoretical_CDF))

(4) Visualize the power-law

Plot the CCDF of the power-law along with a best-fit line made using your estimated scaling exponent. You can use your code from the power-law visualization assignment from earlier, but you will have to add the best fit line. One way to do this is to generate points for the x-axis using [np.linspace](https://docs.scipy.org/doc/numpy/reference/generated/numpy.linspace.html) (<https://docs.scipy.org/doc/numpy/reference/generated/numpy.linspace.html>) or [np.arange](https://docs.scipy.org/doc/numpy/reference/generated/numpy.arange.html) (<https://docs.scipy.org/doc/numpy/reference/generated/numpy.arange.html>) and then calculate what the y-axis values should be using xmin and alpha for a straight line in the log-log CCDF plot. If the fit is good, it should fall right on-top of the empirical data-points on the CCDF plot.


In [7]:

```
# Your visualization code here
bins = np.logspace(0.0, 3.0, num=40)
Y, X = np.histogram(degree, bins=bins, normed=False)
X = [x*np.sqrt(bins[1]) for x in X[:-1]] # find the center point for each bin.
cum = np.cumsum(Y)
ccum = 1-cum/sum(Y)
plt.ylim((0.001, 1))
plt.xlabel('Node Degree, k')
plt.ylabel("CCDF(k)")
plt.loglog(X,ccum, 'o', markersize=3, markerfacecolor='none', label='Empirical data')

x = np.logspace(0.0, 3.0, num=40)
alpha = fit.power_law.alpha
xmin = fit.xmin
y = (x/xmin)**(-alpha+1)
plt.loglog(x, y, '-r', linestyle='--', label='power-law (alpha = 3, xmin = 9)')
plt.legend(loc='upper right', fontsize=8)
```

Out[7]:

<matplotlib.legend.Legend at 0x1b257b85898>



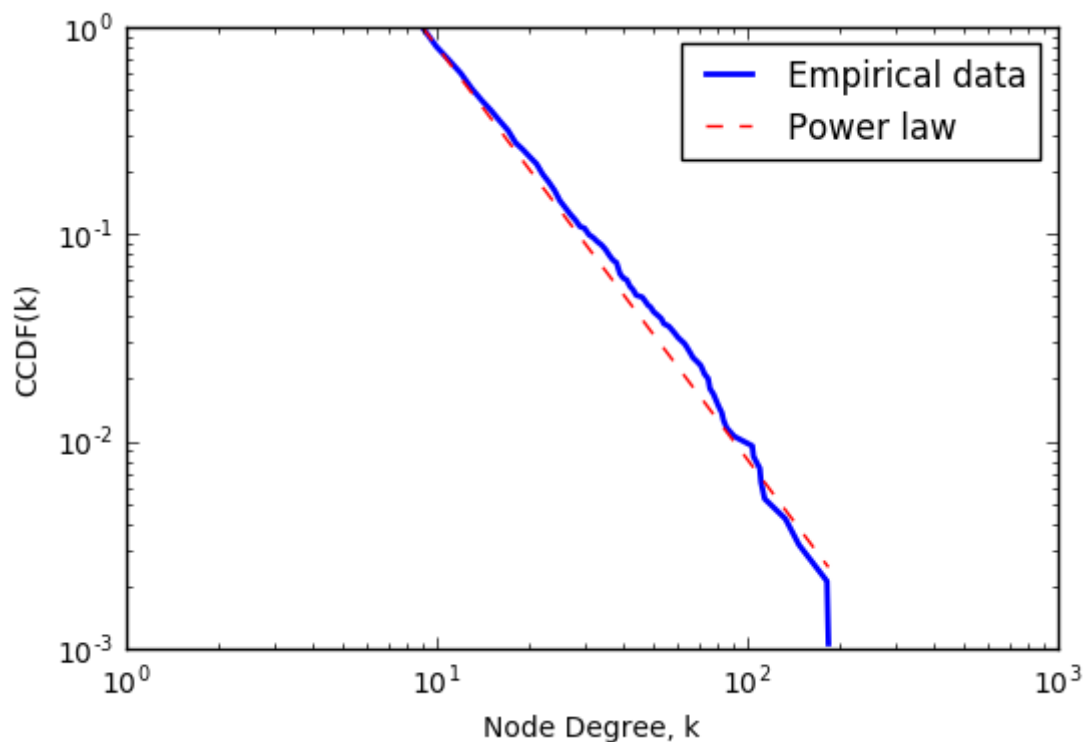
Alternative, we can use powerlaw package's built in function to draw the ccdf plot.

In [8]:

```
fig = fit.plot_ccdf(linewidth=2, label = 'Empirical data')
fit.power_law.plot_ccdf(ax=fig, color='r', linestyle='--', label = 'Power law')
plt.xlabel('Node Degree, k')
plt.ylabel("CCDF(k)")
plt.legend()
```

Out[8]:

<matplotlib.legend.Legend at 0x1b258742b70>



Using your own data

Now that you have done the fit and visualization for the BA graph, find a graph or other data and repeat the analysis. You are welcome to use data from this [graph website \(http://www-personal.umich.edu/~mejn/netdata/\)](http://www-personal.umich.edu/~mejn/netdata/), [Clauset's work \(http://tuvalu.santafe.edu/~aaronc/powerlaws/data.htm\)](http://tuvalu.santafe.edu/~aaronc/powerlaws/data.htm), or from whatever real-world graphs you can find. Make sure to include the graph or data file when uploading this notebook (and avoid using a huge dataset with millions of nodes).

(1) Getting real-world distribution

In [9]:

```
# Your code here
fname = 'fires.txt'
!head 'fires.txt'
nums = [float(x) for x in open(fname)]
print (nums[:5], '...', nums[-5:])

0.10
0.10
0.10
0.10
0.10
0.10
0.10
0.10
0.10
0.10
[0.1, 0.1, 0.1, 0.1, 0.1] ... [177544.0, 178900.0, 187300.0, 400100.0, 41205
0.0]
```

(2) Finding cut-off

In [10]:

```
# Your code here
fit2 = powerlaw.Fit(nums)
print(fit2.xmin)
```

Calculating best minimal value for power law fit
C:\Users\Bao\Anaconda3\lib\site-packages\powerlaw.py:692: RuntimeWarning: in
valid value encountered in true_divide
(Theoretical_CDF * (1 - Theoretical_CDF))
6324.0

(3) Estimating scaling exponent

In [11]:

```
# Your code here
fit2 = powerlaw.Fit(nums, xmin=6324.0)
print(fit2.power_law.alpha)
```

2.16362867855

C:\Users\Bao\Anaconda3\lib\site-packages\powerlaw.py:692: RuntimeWarning: in
valid value encountered in true_divide
(Theoretical_CDF * (1 - Theoretical_CDF))


(4) Visualize power-law and best-fit

In [12]:

```
# Your code here
import pylab
fig2 = fit2.plot_ccdf(linewidth=3, label = 'Empirical data')
fit2.power_law.plot_ccdf(ax=fig2, color='r', linestyle='--', label = 'Power law')
plt.xlabel('Node Degree, k')
plt.ylabel("CCDF(k)")
plt.legend()
```

Out[12]:

<matplotlib.legend.Legend at 0x1b257abfa58>



In []: